

# POMDPStressTesting.jl Example: Walk1D

Robert J. Moss

*Computer Science, Stanford University*

MOSSR@CS.STANFORD.EDU

## Abstract

In this self-contained tutorial, we define a simple problem for adaptive stress testing (AST) to find failures. This problem, called Walk1D, samples random walking distances from a standard normal distribution  $\mathcal{N}(0, 1)$  and defines failures as walking passed a certain threshold (which is set to  $\pm 10$  in this example). AST will select the seed which deterministically controls the sampled value from the distribution (i.e. from the transition model). AST will try to find the set of seeds that maximizes the log-probability of the samples, while simultaneously guiding the simulation to failures using a notion of distance to failure.

## 1. Parameters and Simulation

First, we define the parameters of our simulation.

```
1 @with_kw mutable struct Walk1DParams
2     startx::Float64 = 0 # Starting x-position
3     threshx::Float64 = 10 # +- boundary threshold
4     endtime::Int64 = 30 # Simulate end time
5 end
```

Next, we define a `BlackBox.Simulation` structure.

```
1 # Implement abstract BlackBox.Simulation
2 @with_kw mutable struct Walk1DSim <: BlackBox.Simulation
3     params::Walk1DParams = Walk1DParams() # Parameters
4     x::Float64 = 0 # Current x-position
5     t::Int64 = 0 # Current time
6     distribution::Distribution = Normal(0, 1) # Transition model
7 end
```

## 2. BlackBox.initialize!

Now we override the `BlackBox` interface, starting with the `BlackBox.initialize!` function that initializes the simulation object. Note, each interface function may modify the `sim` object in place.

```
1 function BlackBox.initialize!(sim::Walk1DSim)
2     sim.t = 0
3     sim.x = sim.params.startx
4 end
```

### 3. BlackBox.transition!

Continuing with the BlackBox interface overriding, we define how the simulation will transitions between steps, returning the log-probability.

```
1 function BlackBox.transition!(sim::Walk1DSim)
2     sim.t += 1 # Keep track of time
3     sample = rand(sim.distribution) # Sample value from distribution
4     logprob = logpdf(sim.distribution, sample) # Get log-probability of sample
5     sim.x += sample # Move agent
6     return logprob::Real
7 end
```

### 4. BlackBox.distance!

We define how close we are to a failure event using a non-negative distance metric.

```
1 BlackBox.distance!(sim::Walk1DSim) = max(sim.params.threshx - abs(sim.x), 0)
```

### 5. BlackBox.isevent!

We define an indication that a failure event occurred.

```
1 BlackBox.isevent!(sim::Walk1DSim) = abs(sim.x) >= sim.params.threshx
```

### 6. BlackBox.isterminal!

Similarly, we define an indication that the simulation is in a terminal state.

```
1 BlackBox.isterminal!(sim::Walk1DSim) = BlackBox.isevent!(sim) ||
2     sim.t >= sim.params.endtime
```

### 7. BlackBox.evaluate!

Lastly, we use our already defined interface to evaluate the system under test. Returning the log-probability  $\log(p)$ , distance to an event  $d$ , and event indication  $e$  as output.

```
1 function BlackBox.evaluate!(sim::Walk1DSim)
2     logprob::Real = BlackBox.transition!(sim) # Step simulation
3     distance::Real = BlackBox.distance!(sim) # Calculate miss distance
4     event::Bool = BlackBox.isevent!(sim) # Check event indication
5     return (logprob::Real, distance::Real, event::Bool)
6 end
```

## 8. AST Setup and Running

Setting up our simulation, we instantiate our simulation object and pass that to the Markov decision process (MDP) object of the adaptive stress testing formulation. We use Monte Carlo tree search (MCTS) with progressive widening on the action space. Hyperparameters are passed to the MCTSASTSolver function, which is a simple wrapper around the POMDPs.jl implementation of MCTS. Lastly, we solve the MDP to produce a policy.

```

1 function setup_ast()
2     # Create black-box simulation object
3     sim::BlackBox.Simulation = Walk1DSim()
4
5     # AST MDP formulation object
6     mdp::ASTMDP = ASTMDP(sim)
7     mdp.params.debug = true # record metrics
8     mdp.params.seed = 1
9     mdp.params.top_k = 10
10
11     # Hyperparameters for MCTS-PW as the solver
12     solver = MCTSASTSolver(depth=sim.params.endtime,
13                             exploration_constant=10.0,
14                             k_action=0.1,
15                             alpha_action=0.85,
16                             n_iterations=1000)
17
18     policy = solve(solver, mdp)
19
20     return (policy, mdp, sim)
21 end

```

Now we setup the AST run.

```

1 (policy, mdp, sim) = setup_ast()

```

Then *playout* the policy and output an action trace of the best trajectory.

```

1 action_trace = playout(mdp, policy)

```

We can also *playback* specific trajectories and display intermediate states along the way (in this case, the  $x$  value of the agent).

```

1 final_state = playback(mdp, action_trace, sim->sim.x)

```